



Software

Efficient kernel with DPC++

ZhuWei, IAGS/MLP/TF GPU

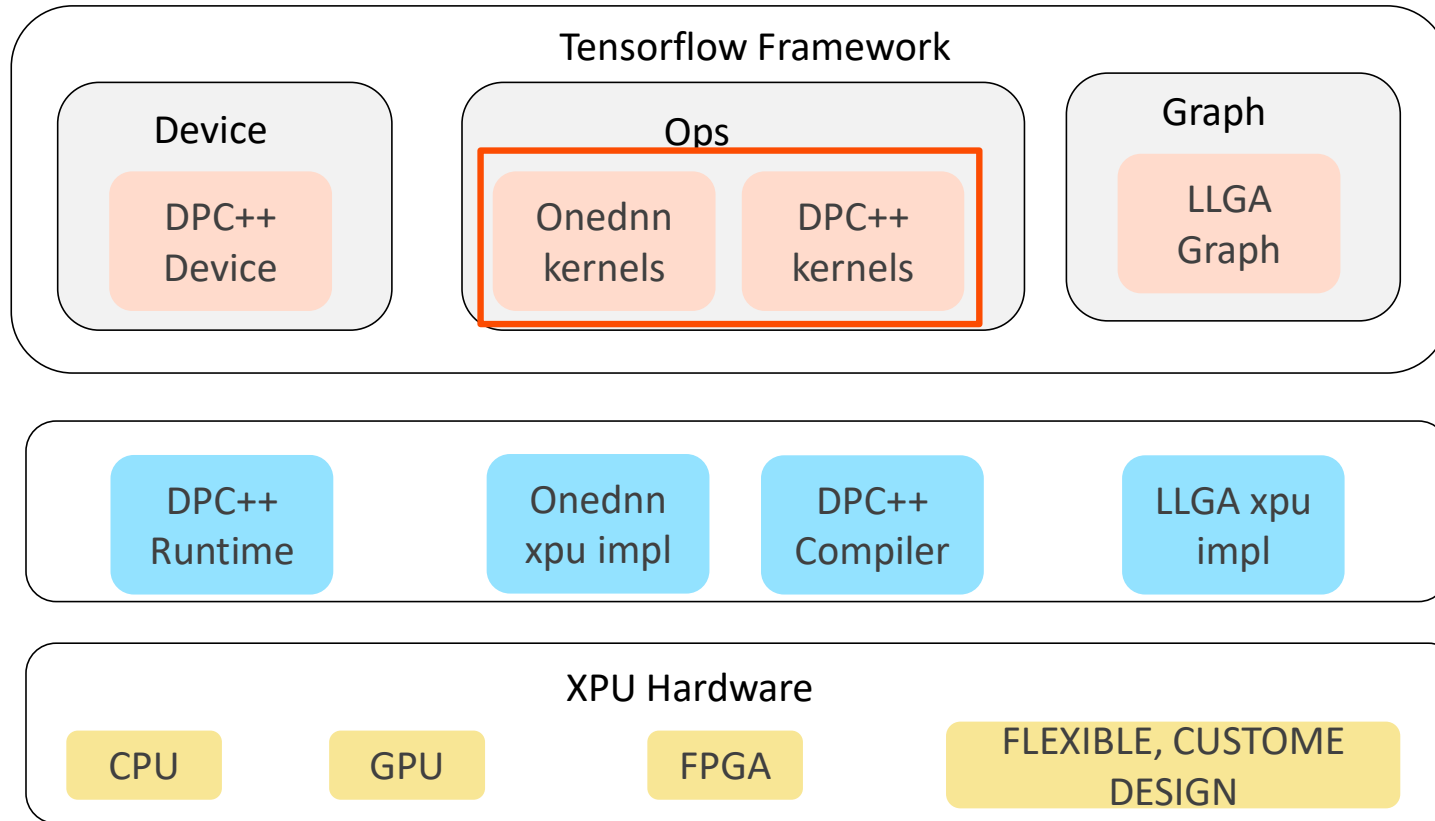
Agenda

- Background
- Efficient Kernel with DPC++
- Q&A

Agenda

- Background
- Efficient Kernel with DPC++
- Q&A

Work Background : TF SoftWare Stack



Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



Work Backgroud

- Functionality

- Op Converage:

- Enabled 49 ops in dev branch, helped to raise op coverage from 64.39% to 91.62%, met 2020 goal.
 - Enabled 20+ ops in Plugin brach.

- Performance(WIP)

- Model analysis: rn50, bert

- Custom Kernel benchmark and Optimization

- Kernel not supported by Onednn(NMS, L2loss, AdamMomentum and so on)
 - Fusion outside oneDNN

Agenda


- Background
- Efficient Kernel with DPC++
 - Migration from CUDA to DPC++
 - Case study : GEMM optimization
- Q&A

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



Migration from CUDA to DPC++: Execution Model



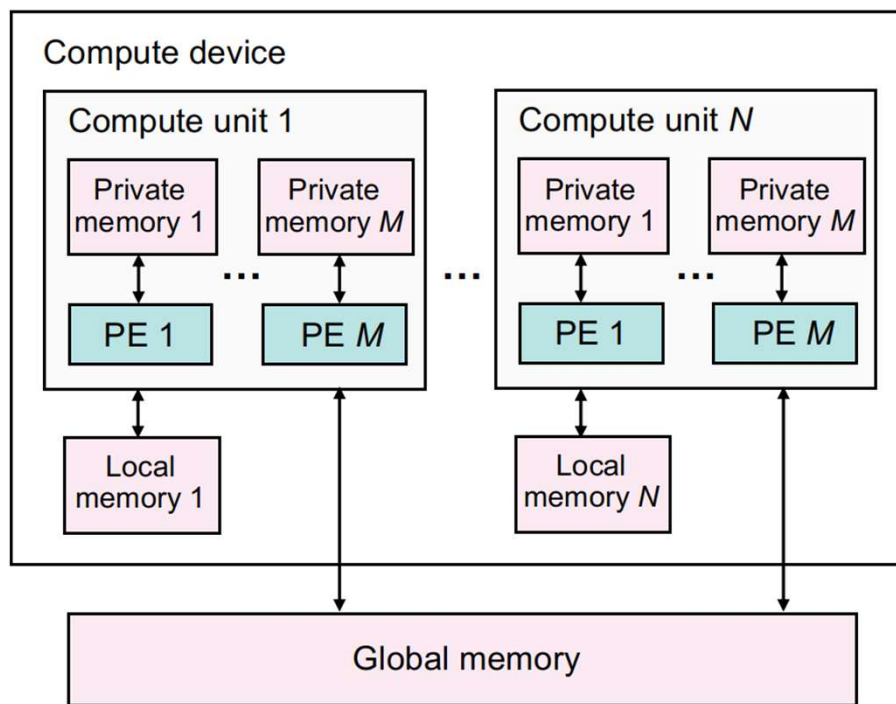
CUDA	DPC++
SP	Process Element
SM	Compute Unit
Thread	Work item
Block	Work group
Grid	NDRange

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



Migration from CUDA to DPC++: Memory Model



CUDA	DPC++
Local memory	Private memory
Shared memory	Local memory
Global memory	Global memory
Constant memory	Constant memory

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



Efficient Kernel with DPC++

- Case study : GEMM optimization

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



What is Our Optimization Goal

- Strive to reach GPU peak performance
- Choose the right metrics:
 - GFLOP/s: for compute-bound kernels
 - Bandwidth: for memory-bound kernels
- Gemm have high arithmetic intensity
- Therefore strive for peak GFLOP/s

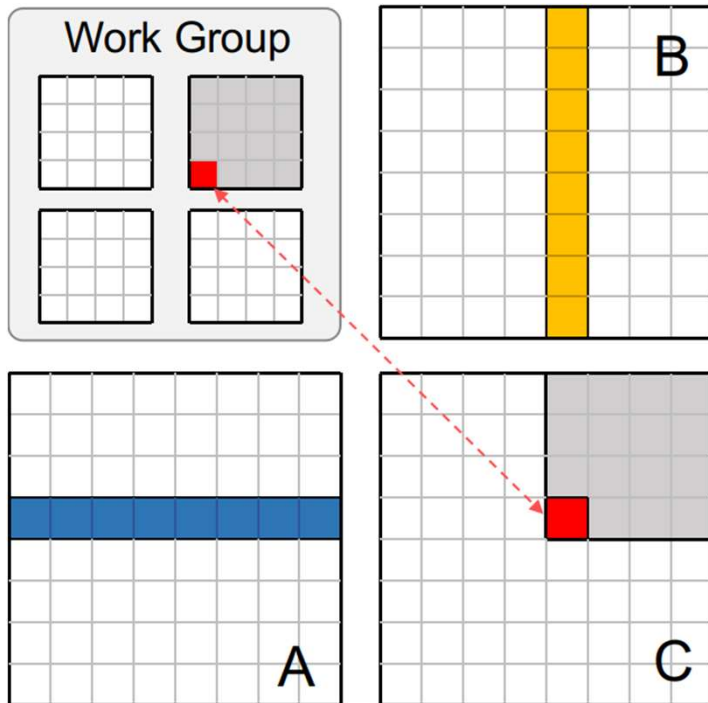
- Use DG1 for this example
 - 96 EUs, 8 threads/EU, 1100 MHz Clock Frequency
 - $96 * 8 * 1.1 * 2 = 1689 \text{ GFLOP/s}$ (Roofline)

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



GEMM #1: Start with DPC++



```
buffer buf_a(A.elements, range<2>{M, N});
buffer buf_b(B.elements, range<2>{N, P});
buffer buf_c(C.elements, range<2>{M, P});
queue q;
q.submit([&](handler &cgh) {
    auto a = buf_a.get_access<access::mode::read>(cgh);
    auto b = buf_b.get_access<access::mode::read>(cgh);
    auto c = buf_c.get_access<access::mode::write>(cgh);

    cgh.parallel_for<MatMulKernel>(range<2>{M, P}, [=](id<2> index) {
        size_t row = index[0];
        size_t col = index[1];
        int val = 0.0;
        for (size_t i = 0; i < N; i += 1) {
            val += a[row][i] * b[i][col];
        }
        c[row][col] = val;
    });
});
q.wait();
```

sycl::malloc_shared
sycl::malloc_device
sycl::aligned_alloc_device

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



Performance

version	Gflops	Time (ms)	Step Speedup	Cumulative Speedup	Efficiency
v1	141.52	29.63	1.0x	1.0x	8%
Peak	1689	2.49	-	-	

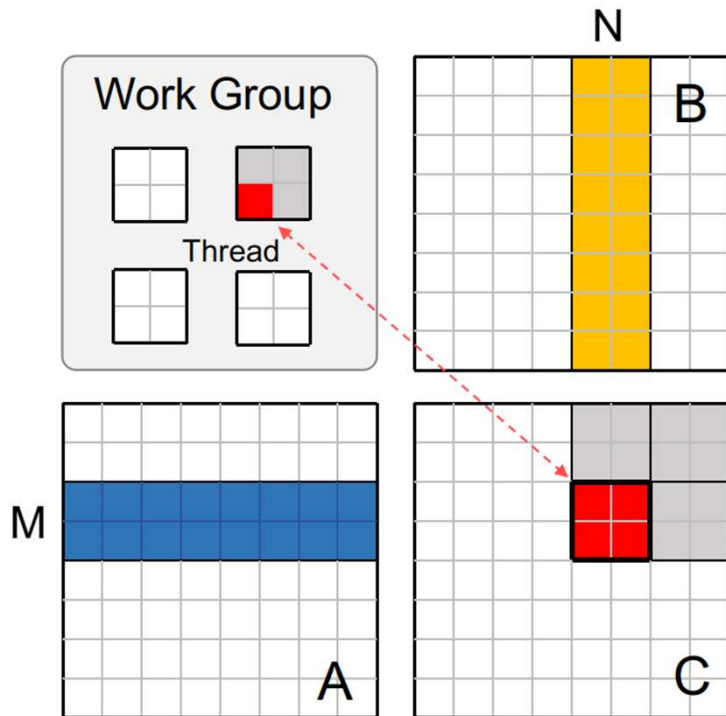
Note: Performance with Matrix size 1280*1280

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



GEMM #2: Multiple Outputs per Thread



Try to increasing cache utilization

```
size_t row = index[0];
size_t col = index[1];
float csub[cm][cn] = {0.0f};
for (int m = 0; m < cm; ++m)
{
    for (int n = 0; n < cn; ++n)
    {
        for (int i = 0; i < N; i += 1)
        {
            csub[m][n] += a[row + m][i] * b[i][col + n];
        }
    }
}
for (int m = 0; m < cm; ++m)
{
    for (int n = 0; n < cn; ++n)
    {
        c[row + m][col + n] += csub[m][n];
    }
}
```

Performance

version	Gflops	Time (ms)	Step Speedup	Cumulative Speedup	Efficiency
v1	141.52	29.63	1.0x	1.0x	8%
v2	71.83	58.39	0.51x	0.51x	4.2%

Why performance drop a lot?

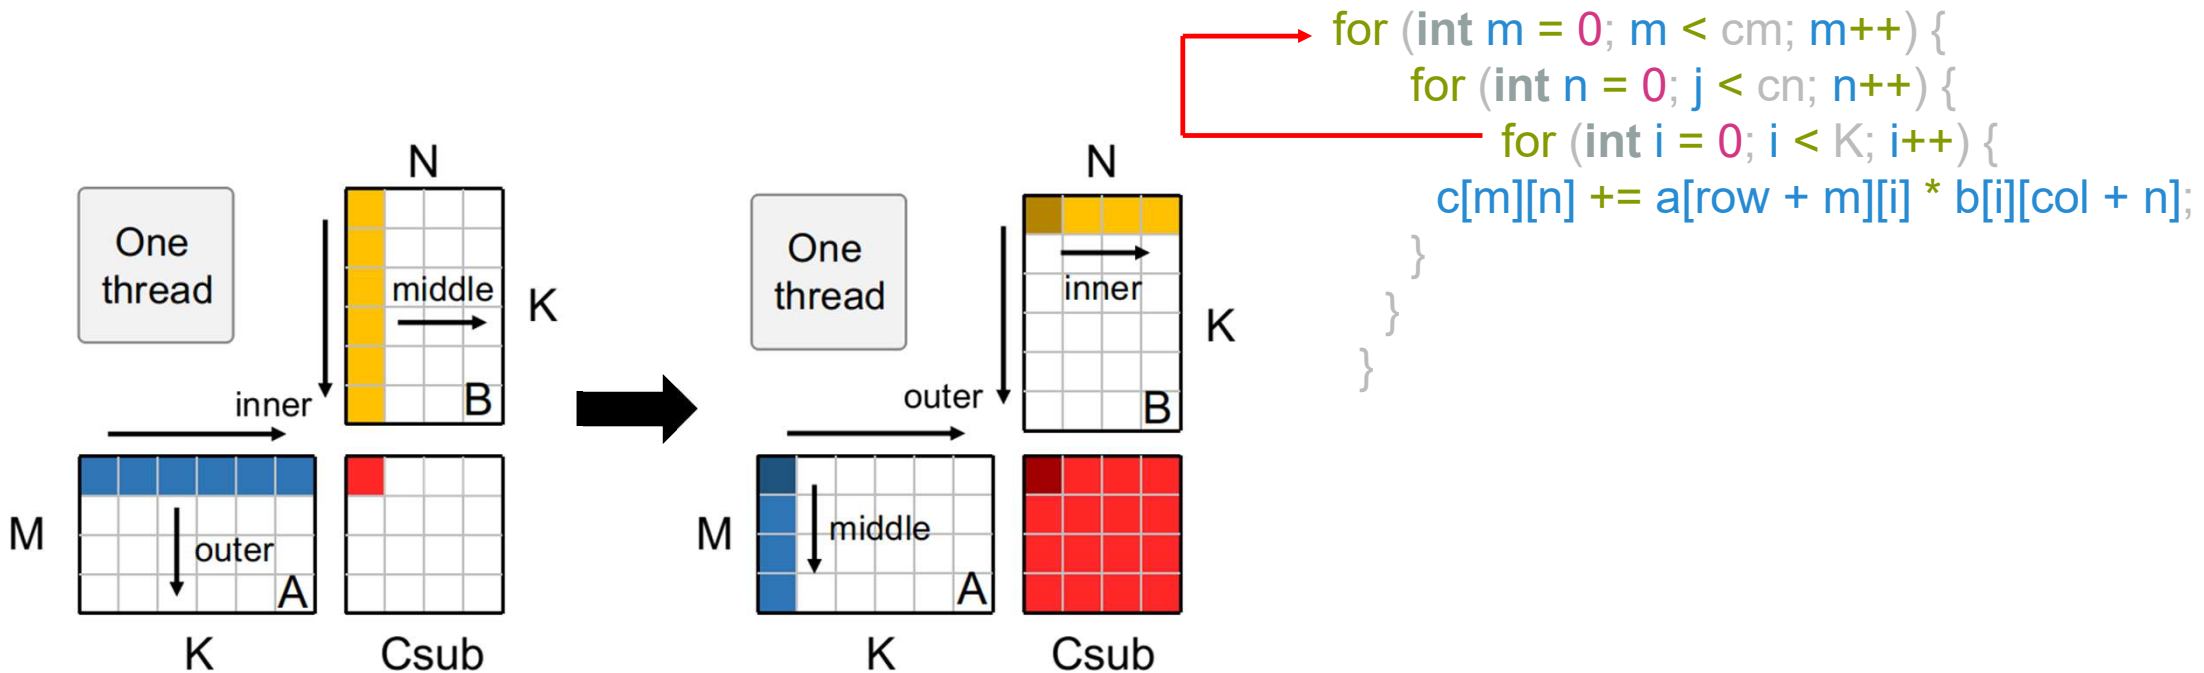
Multiple outputs per thread result in less parallelism and latency can't be hided

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



GEMM #3: Multiple Outputs per Thread -- Permute



- Reduce memory repeat access
- Increase L2 cache hit rate

Performance

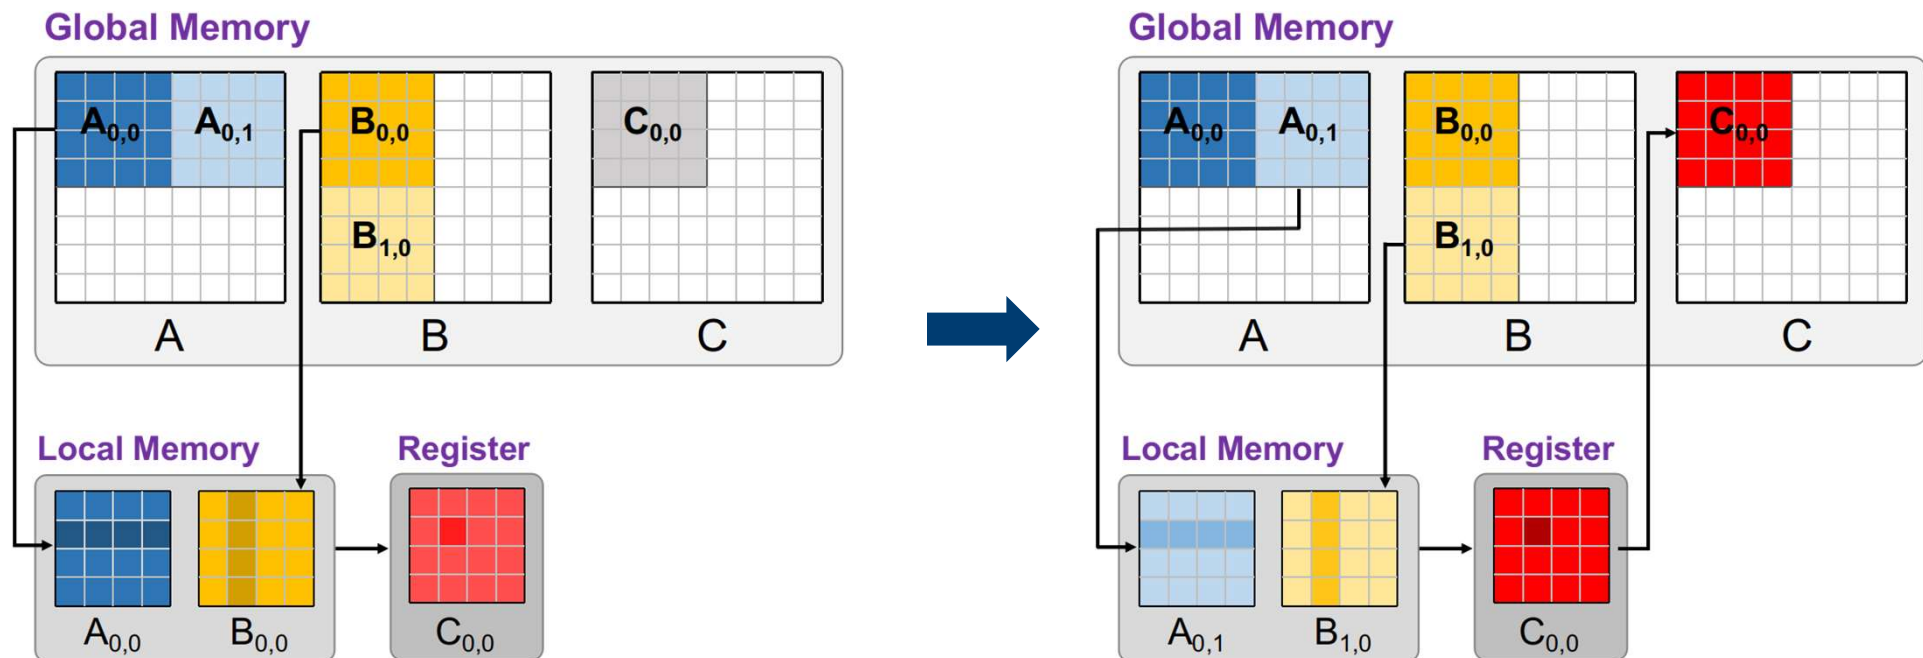
version	Gflops	Time (ms)	Step Speedup	Cumulative Speedup	Efficiency
v1	141.52	29.63	1.0x	1.0x	8.3%
v2	71.83	58.39	0.51x	0.51x	4.2%
v3	609.09	6.88	8.47x	4.31x	36%
Peak	1689	2.49	-	-	

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



GEMM #4: Using Local Memory



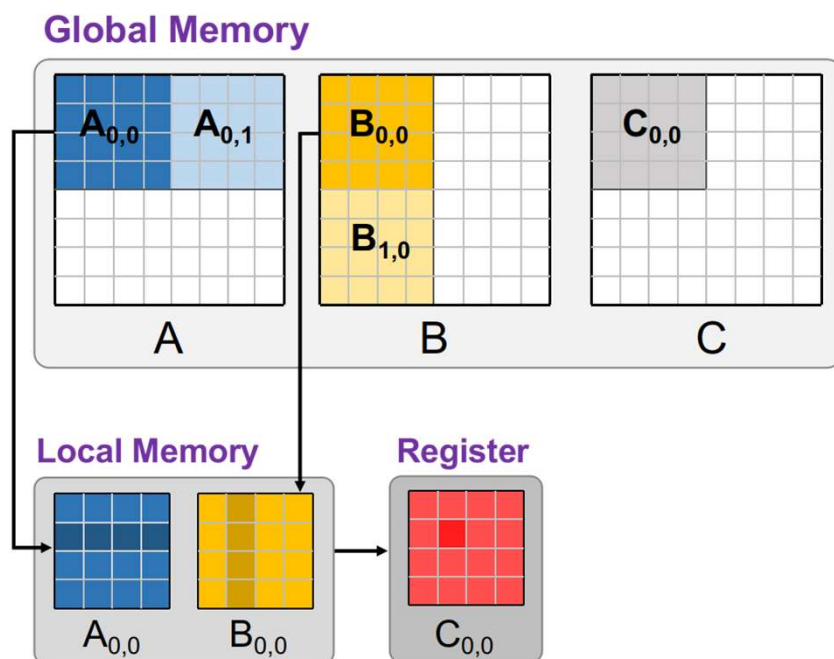
- Local memory is allocated per work group
- User-managed data caches
- Access to Local memory is much faster than global memory

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



GEMM #4: Using Local Memory



```
for (int k = 0; k < N; k += TILE_K) {
    // load A
    for (int m = 0; m < cm; ++m) {
        int tile_row = i_row + m * block_size_row;
        int tile_col = i_col;
        int row = tile_row + g_row * cm * block_size_row;
        int col = tile_col + k;
        Asub[tile_row][tile_col] = A[row * N + col];
    }

    // load B
    for (int p = 0; p < cp; ++p) {
        int tile_row = i_row;
        int tile_col = i_col + p * block_size_col;
        int row = tile_row + k;
        int col = g_col * cp * block_size_col + tile_col;
        Bsub[tile_row][tile_col] = B[row * P + col];
    }

    // wait all memory has been stored to Asub & Bsub
    item.barrier(sycl::access::fence_space::local_space);

    for (int k1 = 0; k1 < TILE_K; ++k1) {
        for (int m = 0; m < cm; ++m) {
            for (int p = 0; p < cp; ++p) {
                Csub[m][p] += Asub[m + i_row * cm][k1] * Bsub[k1][p + i_col * cp];
            }
        }
    }

    item.barrier(sycl::access::fence_space::local_space);
}
```

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



Performance


version	Gflops	Time (ms)	Step Speedup	Cumulative Speedup	Efficiency
v1	141.52	29.63	1.0x	1.0x	8.3%
v2	71.83	58.39	0.51x	0.51x	4.2%
v3	609.09	6.88	8.47x	4.31x	36%
v4	982.14	4.27	1.61x	6.96x	58%
Peak	1689	2.49	-	-	

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



GEMM #5: Local Memory -- Bank Conflicts




```
for (int k = 0; k < N; k += TILE_K) {
    // load A
    for (int m = 0; m < cm; ++m) {
        int tile_row = i_row + m * block_size_row;
        int tile_col = i_col;
        int row = tile_row + g_row * cm * block_size_row;
        int col = tile_col + k;
        Asub[tile_row][tile_col] = A[row * N + col];
    }

    // load B
    for (int p = 0; p < cp; ++p) {
        int tile_row = i_row;
        int tile_col = i_col + p * block_size_col;
        int row = tile_row + k;
        int col = g_col * cp * block_size_col + tile_col;
        Bsub[tile_row][tile_col] = B[row * P + col];
    }

    // wait all memory has been stored to Asub & Bsub
    item.barrier(sycl::access::fence_space::local_space);
}
```

GEMM #5: Local Memory -- Bank Conflicts



- Different threads in one warp access “different” words in the same bank.


```
for (int k = 0; k < N; k += TILE_K) {
    // load A
    for (int m = 0; m < cm; ++m) {
        int tile_row = i_row + m * block_size_row;
        int tile_col = i_col;
        int row = tile_row + g_row * cm * block_size_row;
        int col = tile_col + k;
        Asub[tile_row][tile_col] = A[row * N + col];
    }

    // load B
    for (int p = 0; p < cp; ++p) {
        int tile_row = i_row;
        int tile_col = i_col + p * block_size_col;
        int row = tile_row + k;
        int col = g_col * cp * block_size_col + tile_col;
        Bsub[tile_row][tile_col] = B[row * P + col];
    }

    // wait all memory has been stored to Asub & Bsub
    item.barrier(sycl::access::fence_space::local_space);
}
```

GEMM #5: Local Memory -- Bank Conflicts free

`#define BLOCK_SIZE_X 16` \longrightarrow `#define BLOCK_SIZE_X 32`



Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



Performance

version	Gflops	Time (ms)	Step Speedup	Cumulative Speedup	Efficiency
v1	141.52	29.63	1.0x	1.0x	8.3%
v2	71.83	58.39	0.51x	0.51x	4.2%
v3	609.09	6.88	8.47x	4.31x	36%
v4	982.14	4.27	1.61x	6.96x	58%
v5	1048.73	3.99	1.06x	7.43x	62%
Peak	1689	2.49	-	-	

Optimization Notice

Copyright © 2018, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



What to do Next?

- At v5 1048GFLOP/s, we're 62% Efficiency.
- Next Step
 - Vectorization: using float4
 - Reduce instruction overhead: unroll loops
 - Tuning
- Trying \longleftrightarrow Finding Cause \longleftrightarrow Proving

Conclusion

- Use peak performance metrics to guide optimization
- Understand performance characteristics
 - Latency hiding
 - Bank conflict
- Try to identify the type of bottleneck
 - Using profiling tool: vtune
 - Memory, core computation, or instruction overhead
- Learn more about Hardware
- Look into the assembly code

Thank You!